

DATA SCIENCE WITH R

Class 4 – Data Manipulation in R

Topic 1



Manipulating Data Using Base R



INDEX



Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables

Missing value treatment

Using reshape2() to transpose data

Manipulating Character Strings

Using sqldf

Data Manipulation: Base R

Data Manipulation: Base R

- Sub-setting data
- Selecting specified columns
- Adding new columns
- Reordering data (Ascending/Descending order)
- Group wise operations
- Producing contingency tables

Sub-setting data

Manipulating data: Base R (Sub setting)

- Sub setting: Selecting a sub set of rows across all columns

```
> head(oj[ojs$brand=='tropicana',])
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE	WORKWOM
1	2	tropicana	40	9.018695	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
2	2	tropicana	46	8.723231	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
3	2	tropicana	47	8.253228	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
4	2	tropicana	48	8.987197	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
5	2	tropicana	50	9.093357	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
6	2	tropicana	51	8.877382	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5							
1	0.4638871	2.110122	1.142857	1.92728	0.3769266							
2	0.4638871	2.110122	1.142857	1.92728	0.3769266							
3	0.4638871	2.110122	1.142857	1.92728	0.3769266							
4	0.4638871	2.110122	1.142857	1.92728	0.3769266							
5	0.4638871	2.110122	1.142857	1.92728	0.3769266							
6	0.4638871	2.110122	1.142857	1.92728	0.3769266							

Manipulating data: Base R (Sub setting)

- Can use multiple conditions, | (or), & (and) operator

```
> head(oj[ojs$brand=='tropicana'|ojs$brand=='dominicks',])
  store  brand week  logmove feat price  AGE60  EDUC  ETHNIC  INCOME  HHLARGE
1     2 tropicana  40 9.018695    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
2     2 tropicana  46 8.723231    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
3     2 tropicana  47 8.253228    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
4     2 tropicana  48 8.987197    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
5     2 tropicana  50 9.093357    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
6     2 tropicana  51 8.877382    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
  WORKWOM  HVAL150 SSTRDIST  SSTRVOL CPDIST5  CPWVOL5
1 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
2 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
3 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
4 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
5 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
6 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
> dim(oj[ojs$brand=='tropicana'|ojs$brand=='dominicks',])
[1] 19298    17
```


Manipulating data: Base R (Sub setting)

```
> dim(oj[oj$brand=='tropicana' & oj$feat==0,])
[1] 8045    17
> head(oj[oj$brand=='tropicana' & oj$feat==0,])
  store    brand week  logmove feat price    AGE60    EDUC    ETHNIC    INCOME    HHLARGE
1     2  tropicana  40 9.018695    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
2     2  tropicana  46 8.723231    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
3     2  tropicana  47 8.253228    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
4     2  tropicana  48 8.987197    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
5     2  tropicana  50 9.093357    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
6     2  tropicana  51 8.877382    0  3.87 0.2328647 0.2489349 0.1142799 10.55321 0.1039534
  WORKWOM    HVAL150 SSTRDIST  SSTRVOL  CPDIST5    CPWVOL5
1 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
2 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
3 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
4 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
5 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
6 0.3035853 0.4638871 2.110122 1.142857 1.92728 0.3769266
> |
```

Manipulating data: Base R (Sub setting)

- So, far logical sub-setting is discussed.
- Use which() operator to get the index for specific rows

```
> index<-which(oj$brand=="dominicks")
> head(index)
[1] 221 222 223 224 225 226
> head(oj[index,])
  store brand week  logmove feat price  AGE60  EDUC  ETHNIC  INCOME  HHLARGE  WORKWOM  HVAL150
221    2 dominicks 40  9.264829    1  1.59 0.2328647 0.2489349 0.1142799 10.55321 0.1039534 0.3035853 0.4638871
222    2 dominicks 46  8.987197    0  2.69 0.2328647 0.2489349 0.1142799 10.55321 0.1039534 0.3035853 0.4638871
223    2 dominicks 47  8.831712    1  2.09 0.2328647 0.2489349 0.1142799 10.55321 0.1039534 0.3035853 0.4638871
224    2 dominicks 48  7.965546    0  2.09 0.2328647 0.2489349 0.1142799 10.55321 0.1039534 0.3035853 0.4638871
225    2 dominicks 50  7.377759    0  2.09 0.2328647 0.2489349 0.1142799 10.55321 0.1039534 0.3035853 0.4638871
226    2 dominicks 51 10.140297    1  1.89 0.2328647 0.2489349 0.1142799 10.55321 0.1039534 0.3035853 0.4638871
  SSTRDIST  SSTRVOL  CPDIST5  CPWVOL5
221 2.110122 1.142857 1.92728 0.3769266
222 2.110122 1.142857 1.92728 0.3769266
223 2.110122 1.142857 1.92728 0.3769266
224 2.110122 1.142857 1.92728 0.3769266
225 2.110122 1.142857 1.92728 0.3769266
226 2.110122 1.142857 1.92728 0.3769266
> |
```

Logical vectors Vs. which

- which() removes NA values in the logical vector
- It only returns the indices where the logical vector is TRUE

```
> #Consider vector sales with missing values
> sales<-c(100,200,NA,300,400,NA,500,600,700,NA,1000,1500,NA,NA)
> #subset data using logical operator
> sales[sales>600]
[1]    NA    NA  700    NA 1000 1500    NA    NA
> #subset data using which
> sales[which(sales>600)]
[1]  700 1000 1500
```

Selecting Columns

Manipulating data: Base R (Selecting)

- Selecting a specified set of columns

```
> head(oj[,c("week", "brand")])
  week    brand
1   40 tropicana
2   46 tropicana
3   47 tropicana
4   48 tropicana
5   50 tropicana
6   51 tropicana
> dim(oj[,c("week", "brand")])
[1] 28947      2
> |
```

Manipulating data: Base R

- Selecting + Sub-setting

```
> head(oj[o$brand=='tropicana' & o$feat==0,c("week","store")])
  week store
1   40     2
2   46     2
3   47     2
4   48     2
5   50     2
6   51     2
> dim(oj[o$brand=='tropicana' & o$feat==0,c("week","store")])
[1] 8045     2
```

Adding new columns

Manipulating data: Base R

- Adding new columns

```
> oj$logInc<-log(oj$INCOME)
> head(oj)
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE
1	2	tropicana	40	9.018695	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534
2	2	tropicana	46	8.723231	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534
3	2	tropicana	47	8.253228	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534
4	2	tropicana	48	8.987197	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534
5	2	tropicana	50	9.093357	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534
6	2	tropicana	51	8.877382	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534

	WORKWOM	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5	logInc
1	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266	2.35643
2	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266	2.35643
3	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266	2.35643
4	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266	2.35643
5	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266	2.35643
6	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266	2.35643

```
> |
```


Ordering data

Ordering

- `order()` returns the element order that results in a sorted vector

```
> students<-c("John","Tim","Alice","Zeus")
> students
[1] "John"  "Tim"   "Alice" "Zeus"
> order(students)
[1] 3 1 2 4
> students[order(students)]
[1] "Alice" "John"  "Tim"   "Zeus"
```

- Application: Very useful for sorting dataframes

Manipulating data: Base R

- Ordering data

```
> head(oj[order(oj$week),])
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE
1	2	tropicana	40	9.018695	0	3.87	0.2328647	0.2489349	0.11427995	10.55321	0.1039534
111	2	minute.maid	40	8.407378	0	3.17	0.2328647	0.2489349	0.11427995	10.55321	0.1039534
221	2	dominicks	40	9.264829	1	1.59	0.2328647	0.2489349	0.11427995	10.55321	0.1039534
331	5	tropicana	40	8.680672	0	3.66	0.1173680	0.3212257	0.05387528	10.92237	0.1030916
447	5	minute.maid	40	8.348538	0	2.99	0.1173680	0.3212257	0.05387528	10.92237	0.1030916
563	5	dominicks	40	7.491088	1	1.59	0.1173680	0.3212257	0.05387528	10.92237	0.1030916

	WORKWOM	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5
1	0.3035853	0.4638871	2.110122	1.1428571	1.927280	0.3769266
111	0.3035853	0.4638871	2.110122	1.1428571	1.927280	0.3769266
221	0.3035853	0.4638871	2.110122	1.1428571	1.927280	0.3769266
331	0.4105680	0.5358834	3.801998	0.6818182	1.600573	0.7363068
447	0.4105680	0.5358834	3.801998	0.6818182	1.600573	0.7363068
563	0.4105680	0.5358834	3.801998	0.6818182	1.600573	0.7363068

Manipulating data: Base R

- Ordering data

```
> head(oj[order(-oj$week),])
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE
110	2	tropicana	160	8.669743	0	2.97	0.2328647	0.2489349	0.11427995	10.55321	0.1039534
220	2	minute.maid	160	10.626582	1	2.19	0.2328647	0.2489349	0.11427995	10.55321	0.1039534
330	2	dominicks	160	9.064158	0	1.82	0.2328647	0.2489349	0.11427995	10.55321	0.1039534
446	5	tropicana	160	8.921057	0	2.78	0.1173680	0.3212257	0.05387528	10.92237	0.1030916
562	5	minute.maid	160	10.825840	1	2.19	0.1173680	0.3212257	0.05387528	10.92237	0.1030916
678	5	dominicks	160	8.723231	0	1.85	0.1173680	0.3212257	0.05387528	10.92237	0.1030916
	WORKWOM	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5					
110	0.3035853	0.4638871	2.110122	1.1428571	1.927280	0.3769266					
220	0.3035853	0.4638871	2.110122	1.1428571	1.927280	0.3769266					
330	0.3035853	0.4638871	2.110122	1.1428571	1.927280	0.3769266					
446	0.4105680	0.5358834	3.801998	0.6818182	1.600573	0.7363068					
562	0.4105680	0.5358834	3.801998	0.6818182	1.600573	0.7363068					
678	0.4105680	0.5358834	3.801998	0.6818182	1.600573	0.7363068					

Manipulating data: Base R

- Subsetting data: Using logical subsets and `which()` statement
- Selecting columns: Using column names at column index
- Adding new columns: Use of `$` operator
- Re-ordering data: `order()`
- Group Wise Summaries
- Producing Contingency tables

GroupWise operations

Manipulating data: Base R

- GroupWise operations
- `tapply()`, `aggregate()`
- What is the mean price of each brand of juice across all stores?

```
> aggregate(oj$price,by=list(oj$brand),mean)
  Group.1      x
1 dominicks 1.735809
2 minute.maid 2.241162
3 tropicana 2.870493
> class(aggregate(oj$price,by=list(oj$brand),mean))
[1] "data.frame"
> tapply(oj$price,oj$brand,mean)
 dominicks minute.maid tropicana 
 1.735809   2.241162   2.870493 
> class(tapply(oj$price,oj$brand,mean))
[1] "array"
> |
```

Manipulating data: Base R

- GroupWise operations
- `tapply()`, `aggregate()`
- What is the mean income level corresponding to brand of juice across all stores?

```
> aggregate(oj$INCOME, by=list(oj$brand), mean)
  Group.1      x
1 dominicks 10.61673
2 minute.maid 10.61673
3 tropicana 10.61673
> class(aggregate(oj$INCOME, by=list(oj$brand), mean))
[1] "data.frame"
> tapply(oj$INCOME, oj$brand, mean)
 dominicks minute.maid tropicana 
 10.61673   10.61673   10.61673 
> class(tapply(oj$INCOME, oj$brand, mean))
[1] "array"
```


Contingency tables

Manipulating data: Base R

- Category wise counts: Contingency tables

Income	Age	Gender	Location
10,000,000	24	M	Arizona
20,000,000	32	F	California
15,000,000	28	M	Arizona
18,000,000	26	F	California

Manipulating data: Base R

- Category wise counts: Contingency tables

Counts	California	Arizona
Male	0	2
Female	2	0

Income	California	Arizona
Male	0	10,000,000+15,000,000
Female	20,000,000+18,000,000	

Manipulating data: Base R

- Category wise counts: Contingency tables
- `table()`, `xtab()`
- Number of people who bought different brands categorized by presence of advertising campaigns

```
> table(oj$brand,oj$feat)
```

	0	1
dominicks	7169	2480
minute.maid	6865	2784
tropicana	8045	1604

Manipulating data: Base R

- Category wise counts: Contingency tables
- `table()`, `xtab()`
- Total income categorized by brand and presence of advertisements

```
> xtabs(oj$INCOME~oj$brand+oj$feat)
```

```
      oj$feat  
oj$brand      0      1  
dominicks  76110.24 26330.63  
minute.maid 72887.96 29552.91  
tropicana   85410.46 17030.41
```

RECAP

- Sub-setting data
- Selecting specified columns
- Adding new columns
- Reordering data (Ascending/Descending order)
- Group wise operations
- Producing contingency tables



Class 4 – Data Manipulation in R

Topic 2



Using dplyr to Manipulate Data



INDEX



Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables

Missing value treatment

Using reshape2() to transpose data

Manipulating Character Strings

Using sqldf



Data Manipulation: dplyr



Manipulating data: dplyr

- dplyr: Whats and Whys
- Sub-setting data using filter()
- Selecting columns using select()
- Adding new columns using mutate()
- Ordering data using arrange()
- Summarizing using summarize() and group_by()
- Using functional pipelines to do more than one manipulation task



Manipulating data: dplyr

- Base R: Good for Medium sized data sets, Awkward Syntax
- dplyr: Faster and elegant syntax
- dplyr: Dataframes
- `install.packages("dplyr")`
- `library(dplyr)`



Sub-setting: filter()

Manipulating data: dplyr

- Sub-setting the data using filter(), base R equivalents: logical subsets and which()
- Only that portion of data such that brand bought is “tropicana”

```
> library(dplyr)
> head(filter(oj, brand=="tropicana"))
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE	WORKWOM
1	2	tropicana	40	9.018695	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
2	2	tropicana	46	8.723231	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
3	2	tropicana	47	8.253228	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
4	2	tropicana	48	8.987197	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
5	2	tropicana	50	9.093357	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853
6	2	tropicana	51	8.877382	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853

	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5
1	0.4638871	2.110122	1.142857	1.92728	0.3769266
2	0.4638871	2.110122	1.142857	1.92728	0.3769266
3	0.4638871	2.110122	1.142857	1.92728	0.3769266
4	0.4638871	2.110122	1.142857	1.92728	0.3769266
5	0.4638871	2.110122	1.142857	1.92728	0.3769266
6	0.4638871	2.110122	1.142857	1.92728	0.3769266

Manipulating data: dplyr

- Sub-setting the data using filter(), base R equivalents: logical subsets and which()
- Only that portion of data such that brand bought is “tropicana” or “dominicks”

```
> head(filter(oj,brand=="tropicana"|brand=="dominicks"))
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE	WORKWOM	HVAL150	SSTRDIST
1	2	tropicana	40	9.018695	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853	0.4638871	2.110122
2	2	tropicana	46	8.723231	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853	0.4638871	2.110122
3	2	tropicana	47	8.253228	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853	0.4638871	2.110122
4	2	tropicana	48	8.987197	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853	0.4638871	2.110122
5	2	tropicana	50	9.093357	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853	0.4638871	2.110122
6	2	tropicana	51	8.877382	0	3.87	0.2328647	0.2489349	0.1142799	10.55321	0.1039534	0.3035853	0.4638871	2.110122

	SSTRVOL	CPDIST5	CPWVOL5
1	1.142857	1.92728	0.3769266
2	1.142857	1.92728	0.3769266
3	1.142857	1.92728	0.3769266
4	1.142857	1.92728	0.3769266
5	1.142857	1.92728	0.3769266
6	1.142857	1.92728	0.3769266



Selecting Columns: `select()`

Manipulating data: dplyr

- Selecting columns from data using select(), base R equivalents: index subsets
- Selecting columns brand and income

```
> head(select(oj, brand, INCOME, feat))
```

	brand	INCOME	feat
1	tropicana	10.55321	0
2	tropicana	10.55321	0
3	tropicana	10.55321	0
4	tropicana	10.55321	0
5	tropicana	10.55321	0
6	tropicana	10.55321	0

```
> |
```


Manipulating data: dplyr

- Selecting columns from data using select(), base R equivalents: index subsets
- Dropping columns brand and income

```
> head(select(oj,-brand,-INCOME,-feat))
```

	store	week	logmove	price	AGE60	EDUC	ETHNIC	HHLARGE	WORKWOM	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5
1	2	40	9.018695	3.87	0.2328647	0.2489349	0.1142799	0.1039534	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266
2	2	46	8.723231	3.87	0.2328647	0.2489349	0.1142799	0.1039534	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266
3	2	47	8.253228	3.87	0.2328647	0.2489349	0.1142799	0.1039534	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266
4	2	48	8.987197	3.87	0.2328647	0.2489349	0.1142799	0.1039534	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266
5	2	50	9.093357	3.87	0.2328647	0.2489349	0.1142799	0.1039534	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266
6	2	51	8.877382	3.87	0.2328647	0.2489349	0.1142799	0.1039534	0.3035853	0.4638871	2.110122	1.142857	1.92728	0.3769266



Creating New Columns: mutate()

Manipulating data: dplyr

- Adding columns to data using mutate(),
- Adding a new column, log(income)

```
> dim(oj)
[1] 28947    17
> head(mutate(oj, logIncome=log(INCOME)))#Changes not made in oj but its copy
  store    brand week  logmove feat price    AGE60    EDUC    ETHNIC    INCOME    HHLARGE    WORKWOM
1     2  tropicana  40  9.018695    0  3.87  0.2328647  0.2489349  0.1142799  10.55321  0.1039534  0.3035853
2     2  tropicana  46  8.723231    0  3.87  0.2328647  0.2489349  0.1142799  10.55321  0.1039534  0.3035853
3     2  tropicana  47  8.253228    0  3.87  0.2328647  0.2489349  0.1142799  10.55321  0.1039534  0.3035853
4     2  tropicana  48  8.987197    0  3.87  0.2328647  0.2489349  0.1142799  10.55321  0.1039534  0.3035853
5     2  tropicana  50  9.093357    0  3.87  0.2328647  0.2489349  0.1142799  10.55321  0.1039534  0.3035853
6     2  tropicana  51  8.877382    0  3.87  0.2328647  0.2489349  0.1142799  10.55321  0.1039534  0.3035853
  HVAL150 SSTRDIST  SSTRVOL CPDIST5  CPWVOL5 logIncome
1  0.4638871  2.110122  1.142857  1.92728  0.3769266    2.35643
2  0.4638871  2.110122  1.142857  1.92728  0.3769266    2.35643
3  0.4638871  2.110122  1.142857  1.92728  0.3769266    2.35643
4  0.4638871  2.110122  1.142857  1.92728  0.3769266    2.35643
5  0.4638871  2.110122  1.142857  1.92728  0.3769266    2.35643
6  0.4638871  2.110122  1.142857  1.92728  0.3769266    2.35643
> dim(oj)
[1] 28947    17
```



Ordering data: arrange()

Manipulating data: dplyr

- Ordering data using `order_by()`,
- Order whole data by income in ascending order

```
> head(arrange(oj, INCOME))
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE	WORKWOM
1	75	tropicana	40	8.971067	0	3.87	0.2076995	0.2195485	0.4159995	9.867083	0.06396471	0.3155833
2	75	tropicana	41	8.392990	0	3.87	0.2076995	0.2195485	0.4159995	9.867083	0.06396471	0.3155833
3	75	tropicana	42	9.018695	0	3.87	0.2076995	0.2195485	0.4159995	9.867083	0.06396471	0.3155833
4	75	tropicana	43	8.624791	0	3.87	0.2076995	0.2195485	0.4159995	9.867083	0.06396471	0.3155833
5	75	tropicana	44	8.476371	0	3.87	0.2076995	0.2195485	0.4159995	9.867083	0.06396471	0.3155833
6	75	tropicana	45	8.877382	0	3.87	0.2076995	0.2195485	0.4159995	9.867083	0.06396471	0.3155833
	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5							
1	0.496	7.192667	2.230769	1.375126	0.7031819							
2	0.496	7.192667	2.230769	1.375126	0.7031819							
3	0.496	7.192667	2.230769	1.375126	0.7031819							
4	0.496	7.192667	2.230769	1.375126	0.7031819							
5	0.496	7.192667	2.230769	1.375126	0.7031819							
6	0.496	7.192667	2.230769	1.375126	0.7031819							


Manipulating data: dplyr

- Ordering data using `order_by()`,
- Order whole data by income in descending order

```
> head(arrange(oj, -INCOME)
+ )
```

	store	brand	week	logmove	feat	price	AGE60	EDUC	ETHNIC	INCOME	HHLARGE	WORKWOM
1	62	tropicana	40	9.373819	0	3.87	0.2225343	0.5177603	0.0265109	11.2362	0.1039793	0.3227652
2	62	tropicana	41	9.368369	0	3.87	0.2225343	0.5177603	0.0265109	11.2362	0.1039793	0.3227652
3	62	tropicana	42	9.570529	0	3.87	0.2225343	0.5177603	0.0265109	11.2362	0.1039793	0.3227652
4	62	tropicana	43	9.400630	0	3.87	0.2225343	0.5177603	0.0265109	11.2362	0.1039793	0.3227652
5	62	tropicana	44	9.329367	0	3.87	0.2225343	0.5177603	0.0265109	11.2362	0.1039793	0.3227652
6	62	tropicana	45	9.631154	0	3.87	0.2225343	0.5177603	0.0265109	11.2362	0.1039793	0.3227652

	HVAL150	SSTRDIST	SSTRVOL	CPDIST5	CPWVOL5
1	0.9166995	5.452685	0.7058824	2.18405	0.2017224
2	0.9166995	5.452685	0.7058824	2.18405	0.2017224
3	0.9166995	5.452685	0.7058824	2.18405	0.2017224
4	0.9166995	5.452685	0.7058824	2.18405	0.2017224
5	0.9166995	5.452685	0.7058824	2.18405	0.2017224
6	0.9166995	5.452685	0.7058824	2.18405	0.2017224



Summarizing data: summarize() and group_by()

Manipulating data: dplyr

- Summarizing data using `summarize()` and `group_by()`
- `group_by()` makes grouped table, `summarize()` can take this grouped table and produce summaries for different columns
- Mean level of income and standard deviation of income for each brand of orange juice

```
> gr_brand<-group_by(oj,brand)
> summarize(gr_brand,mean(INCOME),sd(INCOME))
Source: local data frame [3 x 3]
```

	brand	mean(INCOME)	sd(INCOME)
1	dominicks	10.61673	0.2823234
2	minute.maid	10.61673	0.2823234
3	tropicana	10.61673	0.2823234



Functional Pipelines: %>%

Manipulating data: dplyr

- dplyr becomes a powerful tool when combined with %>% (pipe) operator
- Several data manipulation tasks can be accomplished in just one line of code
- Traditionally functional composition is achieved by using nested function calls
- For example, Find the mean price for all people whose income is ≥ 10.5

```
> #Base R code
> mean(oj[oj$INCOME>=10.5,"price"])
[1] 2.270229
> #dplyr code
> summarize(filter(oj, INCOME>=10.5), mean(price))
  mean(price)
1      2.270229
```

Manipulating data: dplyr

```
> oj%>%filter(INCOME>=10.5)%>%summarize(mean(price))
```



filter(oj, INCOME>=10.5)

Manipulating data: dplyr

```
> oj%>%filter(INCOME>=10.5)%>%summarize(mean(price))
```

filter(oj, INCOME>=10.5)

Subsetted data, input for
summarize()

Manipulating data: dplyr

- Clearly the code looks very messy, using a %>% operator, we can make it more readable

```
> oj%>%filter(INCOME>=10.5)%>%summarize(mean(price))
  mean(price)
1      2.270229
```

- This can be easily read as:
- Take data oj, filter it based on income
- Take this filtered data frame and compute the mean of price

Manipulating data: dplyr

- Subset the data based on $\text{price} \geq 2.5$, create a column `logIncome`, compute the mean, standard deviation and median of column `logIncome`

```
> oj %>% filter(price >= 2.5) %>% mutate(logIncome = log(INCOME)) %>% summarize(mean(logIncome), median(logIncome), sd(logIncome))
# A tibble: 1 x 3
  mean(logIncome) median(logIncome) sd(logIncome)
  <dbl>           <dbl>           <dbl>
1 2.360997       2.363903       0.02800802
```



RECAP

- dplyr: better manipulation functionality
- Sub-setting data using `filter()`
- Selecting columns using `select()`
- Adding new columns using `mutate()`
- Ordering data using `arrange()`
- Summarizing using `summarize()` and `group_by()`
- Using functional pipelines to do more than one manipulation task



Class 4 – Data Manipulation in R

Topic 3



Working with Date Objects



INDEX



Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables

Missing value treatment

Using reshape2() to transpose data

Manipulating Character Strings

Using sqldf



Manipulating Date objects



Manipulating date objects

- Dates are treated as a special data type in most programming languages
- R also treats dates as a separate data type
- Doing so, one can easily work with dates
- Usual date operations include:
 - Finding time interval between two points in data: age
 - Extracting months and week days



Manipulating date objects

- Character to date conversion-Date Class
- Extracting months and weekdays
- Using `difftime()`
- Manipulating data involving dates
- POSIXct and POSIXlt Classes
- Working with `lubridate()`

Manipulating date objects

- Using Date class to convert a character into a Date

```
> fd<-read.csv("Fd.csv")
> str(fd)
'data.frame': 30443 obs. of 25 variables:
 $ FlightDate      : Factor w/ 74 levels "01-Feb-14","01-Jan-14",...: 2 2 2 2 2 2 2 2 2 2 ...
 $ UniqueCarrier   : Factor w/ 13 levels "AA","AS","B6",...: 4 4 4 4 11 11 11 11 11 4 ...
 $ AirlineID       : int  19790 19790 19790 19790 20355 20355 20355 20355 20355 19790 ...
 $ Carrier        : Factor w/ 13 levels "AA","AS","B6",...: 4 4 4 4 11 11 11 11 11 4 ...
 $ TailNum        : Factor w/ 2816 levels "", "D942DN", "NOEGMQ",...: 2641 2512 2490 2581 1657 34 60 17
98 1443 473 ...
 $ FlightNum       : int  335 1095 2422 1607 657 894 1843 2041 413 2030 ...
 $ OriginAirportID : int  11057 11057 11057 11057 11057 11057 11057 11057 11057 13232 ...
 $ OriginAirportSeqID: int  1105703 1105703 1105703 1105703 1105703 1105703 1105703 1105703 1105703 132
```

Manipulating date objects

- Using Date class to convert a character into a Date

```
> fd$FlightDate<-as.Date(fd$FlightDate,"%d-%b-%y")
> str(fd)
'data.frame': 30443 obs. of 25 variables:
 $ FlightDate      : Date, format: "2014-01-01" "2014-01-01" "2014-01-01" ...
 $ UniqueCarrier   : Factor w/ 13 levels "AA","AS","B6",...: 4 4 4 4 11 11 11 11 11 4 ...
 $ AirlineID       : int  19790 19790 19790 19790 20355 20355 20355 20355 20355 19790 ...
 $ Carrier        : Factor w/ 13 levels "AA","AS","B6",...: 4 4 4 4 11 11 11 11 11 4 ...
 $ TailNum        : Factor w/ 2816 levels "", "D942DN", "N0EGMQ",...: 2641 2512 2490 2581 1657 34 60 17
98 1443 473 ...
 $ FlightNum       : int  335 1095 2422 1607 657 894 1843 2041 413 2030 ...
 $ OriginAirportID : int  11057 11057 11057 11057 11057 11057 11057 11057 11057 13232 ...
 $ OriginAirportSeqID: int  1105703 1105703 1105703 1105703 1105703 1105703 1105703 1105703 1105703 132
```

Manipulating date objects

- Using Date class to convert a character into a Date

Code	Value
%d	Day of month (decimal number)
%m	Month (decimal number)
%b	Month (abbreviated)
%B	Month (full name)
%y	Year (2 digits)
%Y	Year (4 digits)

- 25/Aug/04: “%d/%b/%y
- 25-August-2004:%d-%B-%Y

Manipulating date objects

- Extracting months and weekdays from data:

```
> head(months(fd$FlightDate))
[1] "January" "January" "January" "January" "January" "January"
> unique(months(fd$FlightDate))
[1] "January" "February" "March"
> head weekdays(fd$FlightDate))
[1] "wednesday" "wednesday" "wednesday" "wednesday" "wednesday" "wednesday"
> unique weekdays(fd$FlightDate))
[1] "Wednesday" "Thursday" "Friday" "Saturday" "Sunday" "Monday" "Tuesday"
```


Manipulating date objects

- Computing time intervals and using `difftime()`

```
> fd$FlightDate[60]-fd$FlightDate[900]
```

```
Time difference of -3 days
```

```
> difftime(fd$FlightDate[3000],fd$FlightDate[90],units = "weeks")
```

```
Time difference of 1.571429 weeks
```

```
> difftime(fd$FlightDate[3000],fd$FlightDate[90],units = "days")
```

```
Time difference of 11 days
```

```
> difftime(fd$FlightDate[3000],fd$FlightDate[90],units = "hours")
```

```
Time difference of 264 hours
```

Manipulating date objects

- Manipulating data involving dates
- Sub-setting data: All rows when the day is Sunday

```
> library(dplyr)
> #Subset the data for day=Sunday
> dim(fd)
[1] 30443    25
> fd_s<-fd%>%filter (weekdays(FlightDate)=="Sunday")
> dim(fd_s)
[1] 4015    25
```

Manipulating date objects

- Manipulating data involving dates
- Find the number of flights on Sundays for destination Atlanta

```
> #Find the number of flights on sundays for destination Atlanta  
> fd%>%filter(weekdays(FlightDate)=="Sunday",DestCityName=="Atlanta, GA")%>%nrow()  
[1] 683
```

Manipulating date objects

- Manipulating data involving dates
- Find the number of flights on Sundays for all cities

```
> #Find the number of flights on Sundays for all cities  
> fd%>%filter(weekdays(FlightDate)=="Sunday")%>%group_by(DestCityName)%>%summarize(n())  
Source: local data frame [10 x 2]
```

	DestCityName	n()
1	Atlanta, GA	683
2	Charlotte, NC	342
3	Chicago, IL	193
4	Denver, CO	448
5	Houston, TX	155
6	Las Vegas, NV	507
7	Los Angeles, CA	603
8	New York, NY	349
9	Phoenix, AZ	466
10	washington, DC	269

Manipulating date objects

- Whenever data has time information along with date, R uses POSIXct and POSIXlt classes to deal with dates

```
> date1<-Sys.time()
> date1
[1] "2015-03-02 17:35:47 IST"
> class(date1)
[1] "POSIXct" "POSIXt"
> weekdays(date1)
[1] "Monday"
> months(date1)
[1] "March"
```

Manipulating date objects

- Whenever data has time information along with date, we use POSIXct and POSIXlt classes to deal with dates

```
> date2<-as.POSIXlt(date1)
> date2
[1] "2015-03-02 17:35:47 IST"
> str(date2)
POSIXlt[1:1], format: "2015-03-02 17:35:47"
```

```
> date2$wday
[1] 1
> date2$zone
[1] "IST"
> date2$hour
[1] 17
> date2$wday
[1] 1
> date2$zone
[1] "IST"
> date2$hour
[1] 17
```

Manipulating date objects

- lubridate() is a package that is a wrapper for POSIXct class
- It has a very simple syntax

```
> library(lubridate)
> fd$FlightDate<-dmy(fd$FlightDate)
> str(fd)
'data.frame': 30443 obs. of 25 variables:
 $ FlightDate      : POSIXct, format: "2014-01-01" "2014-01-01" "2014-01-01" ...
 $ UniqueCarrier   : Factor w/ 13 levels "AA","AS","B6",...: 4 4 4 4 11 11 11 11 11 4 ...
 $ AirlineID       : int  19790 19790 19790 19790 20355 20355 20355 20355 20355 19790 ...
 $ Carrier        : Factor w/ 13 levels "AA","AS","B6",...: 4 4 4 4 11 11 11 11 11 4 ...
 $ TailNum        : Factor w/ 2816 levels "", "D942DN", "N0EGMQ",...: 2641 2512 2490 2581 1657 34 60 17
98 1443 473 ...
 $ FlightNum       : int  335 1095 2422 1607 657 894 1843 2041 413 2030 ...
 $ OriginAirportID : int  11057 11057 11057 11057 11057 11057 11057 11057 11057 13232 ...
 $ OriginAirportSeqID: int  1105703 1105703 1105703 1105703 1105703 1105703 1105703 1105703 1105703 132
3202 ...
 $ OriginCityMarketID: int  31057 31057 31057 31057 31057 31057 31057 31057 31057 30977 ...
 $ Origin          : Factor w/ 10 levels "ATL","CLT","DCA",...: 2 2 2 2 2 2 2 2 2 9 ...
 $ OriginCityName   : Factor w/ 10 levels "Atlanta, GA",...: 2 2 2 2 2 2 2 2 2 3 ...
 $ OriginState      : Factor w/ 10 levels "AZ","CA","CO",...: 6 6 6 6 6 6 6 6 6 5 ...
```

Manipulating date objects

- lubridate() is a package that is a wrapper for POSIXct class
- It has a very simple syntax.

Function	Date
dmy()	26/11/2008
ymd()	2008/11/26
mdy()	11/26/2008
dmy_hm()	26/11/2008 20:15
dmy_hms()	26/11/2008 20:15:30



RECAP

- Character to date conversion-Date Class
- Extracting months and weekdays
- Using `difftime()`
- Manipulating data involving dates
- `POSIXct` and `POSIXlt` Classes
- Working with `lubridate()`



Class 4 – Data Manipulation in R

Topic 4

★ Merging Tables ★

INDEX

Manipulating data using base R

Using dplyr to manipulate data

Working with date objects



Merging tables

Missing value treatment

Using reshape2() to transpose data

Manipulating Character Strings

Using sqldf



Joining dataframes



Joining dataframes

- Just like tables can be joined in sql, we can perform joins on dataframes in R
- Following types of joins can be accomplished
 - Inner Join
 - Left outer join
 - Right outer join
 - Full outer join

Joining dataframes

- Inner join: Joining two tables based on a key column, such that rows matching in both tables are selected

	CustomerId	Product
1	1	Toaster
2	2	Toaster
3	3	Toaster
4	4	Radio
5	5	Radio
6	6	Radio

	CustomerId	State
1	2	Alabama
2	4	Alabama
3	6	Ohio

Joining dataframes

- Inner join: Joining two tables based on a key column, such that rows matching in both tables are selected

	CustomerId	Product
1	1	Toaster
2	2	Toaster
3	3	Toaster
4	4	Radio
5	5	Radio
6	6	Radio

	CustomerId	State
1	2	Alabama
2	4	Alabama
3	6	Ohio

```
> merge(x=df1,y=df2,by="CustomerId")#Inner Join/Intersection of both tables
```

	CustomerId	Product	State
1	2	Toaster	Alabama
2	4	Radio	Alabama
3	6	Radio	Ohio

Joining dataframes

- Full Outer Join: Two tables are joined irrespective of any match between the rows

	CustomerId	Product
1	1	Toaster
2	2	Toaster
3	3	Toaster
4	4	Radio
5	5	Radio
6	6	Radio

	CustomerId	State
1	2	Alabama
2	4	Alabama
3	6	Ohio

```
> merge(x = df1, y = df2, by = "CustomerId", all = TRUE)#Outer join:
```

	CustomerId	Product	State
1	1	Toaster	<NA>
2	2	Toaster	Alabama
3	3	Toaster	<NA>
4	4	Radio	Alabama
5	5	Radio	<NA>
6	6	Radio	Ohio

Joining dataframes

- Left Outer Join: All the rows of left table are retained while matching rows of right table are displayed

	CustomerId	Product		CustomerId	State
1	1	Toaster	1	2	Alabama
2	2	Toaster	2	4	Alabama
3	3	Toaster	3	6	Ohio
4	4	Radio			
5	5	Radio			
6	6	Radio			

```
> merge(x = df1, y = df2, by = "CustomerId", all.x=TRUE)#Left join
```

	CustomerId	Product	State
1	1	Toaster	<NA>
2	2	Toaster	Alabama
3	3	Toaster	<NA>
4	4	Radio	Alabama
5	5	Radio	<NA>
6	6	Radio	Ohio

Joining dataframes

- Right Outer Join: All the rows of right table are retained while matching rows of left table are displayed

	CustomerId	Product
1	1	Toaster
2	2	Toaster
3	3	Toaster
4	4	Radio
5	5	Radio
6	6	Radio

	CustomerId	State
1	2	Alabama
2	4	Alabama
3	6	Ohio

```
> merge(x = df1, y = df2, by = "CustomerId", all.y=TRUE)#Right join
```

	CustomerId	Product	State
1	2	Toaster	Alabama
2	4	Radio	Alabama
3	6	Radio	Ohio



RECAP

- Inner Join
- Left outer join
- Right outer join
- Full outer join



Class 4 – Data Manipulation in R

Topic 5



Missing Value Treatment



INDEX

Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables



Missing value treatment

Using reshape2() to transpose data

Manipulating Character Strings

Using sqldf



Missing Values



Missing Values

- Identifying missing values in a column
- Imputing missing values

Missing Values

- Using `is.na()` to find out the total number of missing values

```
> a<-c(1,2,3,4,5,6,NA,NA,NA,7,8,9)
> is.na(a)
[1] FALSE FALSE FALSE FALSE FALSE TRUE  TRUE  TRUE FALSE FALSE FALSE
> sum(is.na(a))
[1] 3
```


Missing Values

- Using `is.na()` to find out the total number of missing values

```
> air<-airquality
```

```
> head(air)
```

	Ozone	Solar.R	wind	Temp	Month	Day
1	41	190	7.4	67	5	1
2	36	118	8.0	72	5	2
3	12	149	12.6	74	5	3
4	18	313	11.5	62	5	4
5	NA	NA	14.3	56	5	5
6	28	NA	14.9	66	5	6

Missing Values

- Using `is.na()` to find out the total number of missing values

```
> sum(is.na(air$Ozone))
```

```
[1] 37
```

```
> sum(is.na(air$Solar.R))
```

```
[1] 7
```

```
> summary(air)
```

Ozone	Solar.R	wind	Temp	Month	Day
Min. : 1.00	Min. : 7.0	Min. : 1.700	Min. :56.00	Min. :5.000	Min. : 1.0
1st Qu.: 18.00	1st Qu.:115.8	1st Qu.: 7.400	1st Qu.:72.00	1st Qu.:6.000	1st Qu.: 8.0
Median : 31.50	Median :205.0	Median : 9.700	Median :79.00	Median :7.000	Median :16.0
Mean : 42.13	Mean :185.9	Mean : 9.958	Mean :77.88	Mean :6.993	Mean :15.8
3rd Qu.: 63.25	3rd Qu.:258.8	3rd Qu.:11.500	3rd Qu.:85.00	3rd Qu.:8.000	3rd Qu.:23.0
Max. :168.00	Max. :334.0	Max. :20.700	Max. :97.00	Max. :9.000	Max. :31.0
NA's :37	NA's :7				

Missing Values

- Imputing missing values

```
> #Imputing Missing values  
> air$Ozone[is.na(air$Ozone)]<-45  
> summary(air)
```

Ozone	Solar.R	wind	Temp	Month	Day
Min. : 1.00	Min. : 7.0	Min. : 1.700	Min. : 56.00	Min. : 5.000	Min. : 1.0
1st Qu.: 21.00	1st Qu.: 115.8	1st Qu.: 7.400	1st Qu.: 72.00	1st Qu.: 6.000	1st Qu.: 8.0
Median : 45.00	Median : 205.0	Median : 9.700	Median : 79.00	Median : 7.000	Median : 16.0
Mean : 42.82	Mean : 185.9	Mean : 9.958	Mean : 77.88	Mean : 6.993	Mean : 15.8
3rd Qu.: 46.00	3rd Qu.: 258.8	3rd Qu.: 11.500	3rd Qu.: 85.00	3rd Qu.: 8.000	3rd Qu.: 23.0
Max. : 168.00	Max. : 334.0	Max. : 20.700	Max. : 97.00	Max. : 9.000	Max. : 31.0
	NA's : 7				

Missing Values

- Imputing missing values

```
> #Imputing Missing values  
> air$Solar.R[is.na(air$Solar.R)]<-mean(air$Solar.R,na.rm=TRUE)  
> summary(air)
```

Ozone	Solar.R	wind	Temp	Month	Day
Min. : 1.00	Min. : 7.0	Min. : 1.700	Min. :56.00	Min. :5.000	Min. : 1.0
1st Qu.: 21.00	1st Qu.:120.0	1st Qu.: 7.400	1st Qu.:72.00	1st Qu.:6.000	1st Qu.: 8.0
Median : 45.00	Median :194.0	Median : 9.700	Median :79.00	Median :7.000	Median :16.0
Mean : 42.82	Mean :185.9	Mean : 9.958	Mean :77.88	Mean :6.993	Mean :15.8
3rd Qu.: 46.00	3rd Qu.:256.0	3rd Qu.:11.500	3rd Qu.:85.00	3rd Qu.:8.000	3rd Qu.:23.0
Max. :168.00	Max. :334.0	Max. :20.700	Max. :97.00	Max. :9.000	Max. :31.0



RECAP

- Identifying missing values in a column
- Imputing missing values



Class 4 – Data Manipulation in R

Topic 6

★ Transposing data using reshape2 ★

INDEX

Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables

Missing value treatment



Using reshape2() to transpose data

Manipulating Character Strings

Using sqldf



Transposing data: reshape2()



Transposing data

- Understanding wide and long data formats
- Converting data in wide format to long
- Converting data in long format to wide

Transposing data

- Understanding wide and long data formats
- Most structured data is in wide format: Variables are columns and Row labels identify observations

Person	Age	Weight
Sankar	26	70
Aiyar	24	60
Singh	25	65

Transposing data

- Understanding wide and long data formats
- The same data can be represented as follows (long format)

Persons	Variable	Value
Sankar	Age	26
Sankar	Weight	70
Aiyar	Age	24
Aiyar	Weight	60
Singh	Age	25
Singh	Weight	65

Transposing data

- Converting a wide format data to long format: melt()

```
> library(reshape2)
> person<-c("Sankar","Aiyar","Singh")
> age<-c(26,24,25)
> weight<-c(70,60,65)
> wide<-data.frame(person,age,weight)
> wide
```

	person	age	weight
1	Sankar	26	70
2	Aiyar	24	60
3	Singh	25	65

Transposing data

- Converting a wide format data to long format: melt()

```
> melt(wide)
Using person as id variables
  person variable value
1 Sankar      age    26
2  Aiyar      age    24
3  Singh      age    25
4 Sankar     weight   70
5  Aiyar     weight   60
6  Singh     weight   65
> melt(wide,id.vars="person")
  person variable value
1 Sankar      age    26
2  Aiyar      age    24
3  Singh      age    25
4 Sankar     weight   70
5  Aiyar     weight   60
6  Singh     weight   65
```

Transposing data

- Converting a wide format data to long format: melt()

```
> melt(wide,id.vars="person",variable.names="Demographics",value.name ="Demo_value" )
  person variable Demo_value
1 Sankar      age         26
2  Aiyar      age         24
3  Singh      age         25
4 Sankar    weight         70
5  Aiyar    weight         60
6  Singh    weight         65
> melted<-melt(wide,id.vars="person",variable.names="Demographics",value.name ="Demo_value" )
```

Transposing data

- Converting a long format data to wide format: dcast()

```
> dcast(melted, person~variable, value.var = "Demo_Value")
```

	person	age	weight
1	Aiyar	24	60
2	Sankar	26	70
3	Singh	25	65



RECAP

- Understanding wide and long data formats
- Converting data in wide format to long
- Converting data in long format to wide



Class 4 – Data Manipulation in R

Topic 7

★ Manipulating Character Strings and ★ Using sqldf

INDEX

Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables

Missing value treatment

Using reshape2() to transpose data



Manipulating Character Strings

Using sqldf



Manipulating Strings in R



Manipulating Strings in R

- We'll cover the following R functions used to manipulate character data
- `substr()`
- `nchar()`
- `tolower()`, `toupper()`
- `strsplit()`, `paste()`
- `grep()`, `grepl()`,
- `sub()`, `gsub()`

Manipulating Strings in R

- substr(), nchar(), tolower(), toupper()

```
> a<-"Batman"  
> substr(a,start=2,stop=6)  
[1] "atman"  
> nchar(a)  
[1] 6  
> tolower(a)  
[1] "batman"  
> toupper(a)  
[1] "BATMAN"
```

Manipulating Strings in R

- `strsplit()`, `paste()`, `grep()`, `grepl()`, `sub()`, `gsub()`

```
> b<-"Bat-Man"
> c<-"Bat/Man"
> strsplit(c,split="/")
[[1]]
[1] "Bat" "Man"
```

```
> paste(b,split=c)
[1] "Bat-Man Bat/Man"
> grep("-",b)
[1] 1
> grepl("/ ",c)
[1] TRUE
> sub("-", "/", b)
[1] "Bat/Man"
> d<-"Bat-Ma-n"
> sub("-", "/", d)
[1] "Bat/Ma-n"
> gsub("-", "/", d)
[1] "Bat/Ma/n"
```

INDEX

Manipulating data using base R

Using dplyr to manipulate data

Working with date objects

Merging tables

Missing value treatment

Using reshape2() to transpose data

Manipulating Character Strings



Using sqldf



Using sql commands within R



SQL queries within R: sqldf

- One can easily use sql queries within R using a package called sqldf
- This package is available at CRAN
- `install.package(sqldf)`
- `library(sqldf)`

Manipulating Strings in R

- Using select statement

```
> library(sqldf)
> #Using SELECT statement
> oj_s<-sqldf("select brand, income, feat from oj ")
```

- Using where statement

```
> #Subsetting using where statement
> oi_s<-sqldf("select brand, income, feat from oi where price<3.8 and income<10")
```

- Using order by statement

```
> #Order by statement
> oj_s<-sqldf("select store,brand,week,logmove,feat,price, income from oj order by income asc")
```

SQL queries within R: sqldf

- Using sql functions

```
> #Demo sql functions
> sqldf("select avg(income) from oj")
  avg(income)
1    10.61673
> sqldf("select min(price) from oj")
  min(price)
1      0.52
```



RECAP

- Basic string manipulation functions: tolower, toupper, grep, grepl, sub, gsub
- Executing sql queries in R, using sqldf() package